

S L O P C H E C K E R

SlopChecker

An open-source utility for funding orgs to screen proposals — **AI-generation signals, citation integrity, quote-anchored evidence**. A human decides; the tool never auto-rejects.

107

PRS MERGED

40

ISSUES CLOSED

5

PEOPLE

1

DAY

Nick Wagner

@nawagner

Emerson Brooking

@etbrooking

Dan Parshall

@danparshall

Dominique Ramsawak

@dvanderram

Alexander Varisco

@990991A

The 10× cascade

> why funders need this now



10×

Easier to write a proposal

LLMs draft, polish, and pad grant applications in minutes — what used to take days now takes an afternoon.



10×

More submissions arrive

Funders' inboxes fill with plausible-looking applications — some real, some hallucinated citations wrapped in fluent prose.



10×

Easier to review them

SlopChecker screens each PDF in ~40s: resolves every DOI, quote-checks every claim, flags AI-generation signals — before a human opens it.

AI made submission cheap. **Review has to get cheap too** — or funders stop reading.

The pitch — live at slop-checker.com


> upload a proposal, get one page of evidence

One question, one answer.

"Check whether a grant proposal holds up, before you spend an hour reading it."

Every finding is tied to the sentence it came from. A score is a signal, not a verdict — the tool **never auto-rejects**.

→ [open source](#) · [Cloudflare](#) + [Railway](#)



The evidence report

> every finding, quote-anchored, on one page

Findings sit beside the passage they came from.

Real DOI resolves. Fabricated one doesn't.
Cited quote isn't in the source paper. Pangram score: **0.96**.

Under the rail: an "All checks" ledger — every registered check runs, gaps report themselves.
No LLM prose in the evidence layer.

→ 10 checks · 41s · ships as a paginated PDF

Checks appear beside the passages they refer to — hover either side to see what connects, click to expand/collapse. Red = failed a check · green = passed · indigo = detector score · gray = could not run.

Resilient Information Ecosystems Initiative

Response to OSF-2026-CIV-014 (excerpt)

The information environment confronting democratic institutions has degraded at a pace that outstrips the capacity of existing civil-society infrastructure to respond. Recent scholarship demonstrates that coordinated inauthentic behavior now shifts measurable public-opinion outcomes in as little as seventy-two hours [1]. As Starbird and colleagues observe, "the velocity of synthetic narratives has rendered traditional fact-checking architectures functionally obsolete" [2]. Our initiative responds directly to this crisis.

Our approach leverages a holistic, multi-stakeholder framework designed to foster synergies across the information resilience landscape. By convening diverse voices and harnessing cutting-edge methodologies, we will unlock transformative capacity at the intersection of technology and civil society. The program builds on findings that prebunking interventions achieve durable attitudinal inoculation across all demographic cohorts [3], and that media-literacy training reduces susceptibility to manipulated content by up to 64 percent [4].

Meridian will deliver twelve regional trainings, a 40-country monitoring network, a peer-reviewed evaluation study, and an open-source detection toolkit within the first grant year, at a total cost of \$148,000.

[1] Chen, R. & Okafor, D. (2024). Measuring opinion cascades in networked publics. J. Online Trust & Safety, 3(2). doi:10.54501/jots.v3i2.118

[2] Starbird, K., West, J., & Moran, R. (2023). Synthetic velocity and the collapse of verification. PNAS. doi:10.1073/pnas.221998120

[3] Vraga, E. & Bode, L. (2025). Inoculation at scale. Misinformation Review, 6(1). doi:10.37016/mr-2025-0163

[4] Guess, A. et al. (2020). A digital media literacy intervention increases discernment. PNAS, 117(27). doi:10.1073/pnas.1920498117

C1 Citation [1]

DOI resolves **YES**
Metadata match **YES**

Q1 Quote → citation [2]

DOI resolves **YES**
Metadata match **YES**
Quote found in source **NO**
Paper is real. Quoted sentence does not appear in its full text.

D1 AI detection (Pangram)

Span score **0.98**
Document score **0.96**
Score, not a verdict.

C3 Citation [3]

DOI resolves **NO**
Found in OpenAlex **NO**
No such article. Journal exists; record does not.

C4 Claimed number → citation [4]

DOI resolves **YES**
Metadata match **YES**
Number found in source **NO**
Source reports 26.5%.

ALL CHECKS

CHECK	RESULT	DETAIL
Within page limit	YES	14 / 15 pp
Required sections present	NO	missing: Budget Justification
All DOIs resolve	NO	3 / 4 — ref [3] not found
Citation metadata match	YES	3 / 3 resolving DOIs
Quotes found in sources	NO	0 / 1 — Q1
Claimed numbers found in sources	NO	0 / 1 — C4
AI detection (Pangram)	0.96	document score
Similar to prior submission	0.87	vs. submission #0083

Similar to prior submission **0.87** vs. submission #0083

